


Actor Critic Methods: From Paper to Code

Policy Approximation and its Advantages


What's a Policy?




Policy \rightarrow probability distribution



M.C. \rightarrow deterministic / epsilon soft




Policy is a function of V or Q



Approximate policy directly

Policy Approximation




Parameterize with N.N. weights θ Gradient ascent in performance $J(\theta)$

$$\theta_{t+1} = \theta_t + \alpha \nabla_{\theta} J(\theta_t)$$


Policy Gradient Methods

Advantages of Policy Approximation



Policy is continuous and finite

Works for continuous action spaces



All actions sampled so no dilemma

Action Selection in Policy Gradients

$h(s, a, \theta) \rightarrow$ Numerical preference for state action pair

$$\pi(a|s, \theta) = \frac{\exp(h(s, a, \theta))}{\sum_b \exp(h(s, b, \theta))}$$

Computed by N.N.

Add up to 1



Preference for profitable actions



Accurate estimates

Limitations of (Some) Policy Gradient Methods



Sample inefficiency



Credit assignment problem




Unstable solutions



May still need to learn V , Q

Policy Gradient Theorem

$J(\theta) \stackrel{\text{def}}{=} v_{\pi_\theta}(s_0) \rightarrow$ value of the episode start state



Are the updates helping?

$\nabla J(\theta)$ Depends on distribution of states encountered

Policy Gradient Theorem

$$\nabla J(\theta) \propto \sum_s \mu(s) \sum_a q_\pi(s, a) \nabla_\theta \pi(a|s, \theta)$$

Turns into expectation value

Can learn from experience to improve performance

Conclusion

- Approximate policy with deep N.N.
- Stochastic policy solves explore exploit dilemma
- Handles continuous action spaces
- Policy gradient theorem lets us use experience

